



Transforming Marketplace Trust & Safety Investigations with Multi-Agent AI

Enterprise AI-native operating system for Trust & Safety investigations.

Prepared by

Aniruddha Vanshiv


Trust & Safety | AI Strategy | Enterprise Operations

[Portfolio](#) | [LinkedIn](#)

Executive Problem Statement

Digital marketplaces process millions of transactions every day, yet Trust & Safety investigations continue to rely heavily on fragmented manual workflows. Analysts must collect evidence from multiple systems, interpret **complex policy documents**, inspect product images, evaluate seller behaviour, and document enforcement decisions. As marketplaces grow, these **fragmented workflows** become slower, more expensive, and increasingly **difficult to scale** while maintaining **consistency and compliance**.


Current Investigation Ecosystem



Marketplace Trust & Safety investigations remain fragmented, labour-intensive, and difficult to scale, creating inconsistent decisions and rising operational costs.

Enterprise Challenge


Modern Trust & Safety operations are slowed by fragmented workflows, inconsistent judgment, and rising review costs.



The Investigation Bottleneck

Enterprise marketplaces may process millions of transactions annually, requiring Trust & Safety teams to investigate thousands of high-risk cases while balancing speed, consistency, and compliance.

Hidden Complexity Behind Every Investigation



Why Existing Solutions Fail




Introducing TrustOS

Enterprise AI-native operating system for scalable, explainable Trust & Safety investigations. TrustOS uses multi-agent AI orchestration to coordinate evidence gathering, policy reasoning, image inspection, and decision support in one enterprise workflow.


TrustOS Solution Architecture

How TrustOS orchestrates end-to-end Trust & Safety investigations using specialised AI agents and human oversight.



TrustOS orchestrates specialised AI agents to deliver faster, consistent and explainable investigations while keeping humans accountable for final decisions.

Traditional vs AI-Native Investigation Workflow



Business Value Delivered by TrustOS

Faster Investigations

Consistent Decisions

Explainable AI

Audit-Ready

Expected Enterprise Outcomes

40–60%
Faster Resolution

30–50%
Decision Consistency

25–40%
Higher Throughput

100%
Traceable Decisions

*Illustrative enterprise targets. Actual outcomes depend on implementation maturity, policy complexity, and operating environment.

Responsible AI & Governance

Explainable AI
Evidence-backed recommendations

Human Oversight
Final decisions remain with reviewers

Audit Trails
Complete decision traceability

Policy Compliance
Consistent governance enforcement

Explore TrustOS

Experience the live prototype and explore the complete implementation.

Live Demo

GitHub Repository